A raw table of U.S. newspaper opinion articles from ProQuest US Newsstream, 1980–2024

Data descriptor, version 1.1 (2026-11-08). Companion to the cleaned dataset derived from this table.

Abstract

The table holds 1,226,856 records exported from ProQuest US Newsstream for ten U.S. newspapers with the highest circulation, published 1980–2024, with no ProQuest id repeated. Each record keeps every metadata field of the export, the article text, parsed dates, years and identifiers with explicit quality flags, its file of origin, and a hash and word count of the text. Apart from repeated records and file footers, nothing was removed or corrected, so the table documents the source as delivered, including news items and publications outside the query. A companion table assigns every record one exclusion reason. This descriptor reports how the data were collected and parsed, what they contain, and where they depart from the intended design.

1. Background and summary

The table is the first stage of a corpus built to study how economists communicate with the public through newspaper opinion pieces. It is a faithful reorganisation of the ProQuest exports, kept raw so that other researchers can derive their own samples from it.

Records	1,226,856, no ProQuest id repeated
Removed from the parsed table	145,118 repeated ProQuest ids, 2,991 repeats of the same text under another id and 2,801 file footers (section 2.4)
Export files	2,801
Publication years	1980–2024
Newspapers queried	10
Distinct ProQuest ids	1,225,974
Records with a named author	597,047 (48.7%)

2. Methods

2.1 Collection

- **Source:** ProQuest US Newsstream, through the Universidad de los Andes library proxy.
- **Period:** exports dated 2024-11-22 to 2025-01-26.
- **Newspapers:** 10, selected for having the highest circulation in the U.S. [1]. Online editions are included on purpose, since much current opinion is published online. New York Post is not provided by the subscription.
- **Query:** exact publication name AND `noft(opinion OR op-ed OR editorial)`, a full-text search for those words.
- **Filters:** document type Commentary, Editorial or News; English; date range resumed from the last download.
- **How:** 100 results per page, select all, export every 5 pages (500 records) to a `.txt` file. CAPTCHAs solved by hand.

2.2 Parsing

- Each export file is split at lines of exactly 60 underscores. The last separator is followed by a footer (contact and copyright lines), which is not a record.
- A record is a header, the body (after `Texto completo:`) and a block of `Label: value` metadata lines. The block is read backwards from its last line, one blank-separated block at a time, for as long as the block starts with a known label. This is needed because the optional `Materia` line makes the block start vary, and body lines can look like labels.
- Stray spaces that ProQuest inserts inside metadata values (`En glish`, `Aug 29, 202 2`) are tolerated when parsing. Values are stored as exported; space-free, case-folded copies are stored in the `*_key` columns. The body is never altered.
- Dates, years and identifiers are parsed with an explicit status (`ok`, `missing`, `unparseable`, and `year_only` for dates). No value is silently dropped.
- Every metadata label of the export gets its own column, named in English (section 3.1); fields that are empty in a record stay empty. The header above the text (document link, abstract, links) is kept too.
- Provenance is kept in `source_file` and `record_index`. The article text is stored as exported, together with its SHA-256 hash, word count and character count.

2.3 Quality control

- Per file, the records found are compared with the separator lines counted independently: all 2,801 files contain one footer and as many records as separators.
- The detection of the `Materia` line was cross-checked against a plain text search on one in ten files, with no disagreement (details in `docs/decisions.md`). The cross-check found and fixed a parsing gap that had hidden `Materia` in 3.9% of records.
- Tests use synthetic records only, and every rule is logged with its evidence in `docs/decisions.md`.

2.4 Changes applied to the export

Change	Rows or fields affected	Reason
Each file split into records, one row per record	all records	A file holds hundreds of records
One column per metadata label, with English names, and the header fields kept	all records	Every field that exists is kept
Dates, years and ids parsed, with a status column; space-free `*_key` copies of categories	columns added; exported values unchanged	Comparing categories despite stray spaces
Hash, word count and character count of the text	columns added	Integrity checks
File footers removed	2,801 rows	Not records
Repeated ProQuest ids removed; the copy with the longest body is kept, the earliest download on ties	145,118 rows	Overlapping downloads (section 5.2)
Repeated text removed: same body, title, date and queried newspaper under another id, titles compared without spaces or case; the copy kept is the earliest that is a signed opinion piece, or the earliest of all	2,991 rows	The same article downloaded under two ids (section 5.2)
Nothing else removed or edited	news items, publications outside the query	Left for the cleaned dataset (section 3.2)

3. Data records

3.1 Raw table

One row per record, as `econ_corpus_raw.csv` (UTF-8, comma-separated). It is written from `data/interim/records.parquet`, which also holds the file footers and the repeated ids. The text is ProQuest-licensed and the license allows sharing it for internal use only, so the table is distributed among the project's researchers and is not published. Spreadsheet programs truncate cells above 32,767 characters; read the CSV with a data-analysis library instead.

Column	Type	Content
`source_file`	string	Export file the record comes from
`record_index`	integer	Position in the file, from 0; the footer is last
`is_footer`	boolean	Row is a file footer; false in every row of the shared table
`title`, `author`, `publication`, `section`	string	Metadata as exported; multi-line values joined
`language`, `source_type`, `doc_type`	string	Metadata as exported
`year_raw`, `date_raw`, `id_raw`	string	Publication year, publication date and ProQuest id as exported
`subject`, `location`, `people`, `organization`, `business_index_term`, `keyword`, `classification`	string	Indexing fields (labels `Materia`, `Lugar`, `Personas`, `Empresa/organización`, `Término de indexación de negocios`, `Identificador / palabra clave`, `Clasificación`)
`origin_place_date`, `pages`, `page_count`, `first_page`, `issue`, `volume`, `issue_year`, `column_name`	string	Print location fields (`Lugar y fecha de origen`, `Páginas`, `Número de páginas`, `Primera página`, `Número`, `Tomo`, `Año`, `columna`)
`issn`, `eissn`, `coden`, `publisher`, `publication_place`, `publication_country`, `publication_subject`, `publication_history`	string	Publication fields (`ISSN`, `e-ISSN`, `CODEN`, `Editorial`, `Lugar de publicación`, `País de publicación`, `Materia de publicación`, `Historial de publicaciones`); `publication_history` is empty in every record of this export
`document_url`, `url`, `header_url`, `links`	string	Links: `URL del documento`, `URL`, the document link above the text, and `Enlaces`
`abstract`	string	`Resumen`; usually "Ninguno disponible."
`company`, `corporate_author`, `product_name`, `ethnicity`, `clip_duration`	string	Other fields as exported (`Empresa`, `Autor corporativo/institucional`, `Nombre del producto`, `Etnicidad`, `Duración del clip`)
`copyright`, `last_updated`, `database`, `fulltext_availability`, `accession_number`, `milestone_date`, `document_features`, `source_details`	string	Record-keeping fields (`Copyright`, `Última actualización`, `Base de datos`, `Disponibilidad de texto completo`, `Número de acceso`, `Fecha del hito`, `Características del documento`, `Detalles de fuente`)
`pub_year`, `date_year`, `doc_id`	integer	Parsed year, year of the parsed date, and ProQuest id
`pub_date`	date	Parsed publication date; empty if only a year is given
`year_status`, `date_status`, `id_status`	string	`ok`, `missing`, `unparseable`; dates also `year_only`
`publication_key`, `language_key`, `source_type_key`, `doc_type_key`	string	Space-free, case-folded copies for comparing categories
`has_materia`	boolean	Record has a `Materia` (subject) line
`n_meta_fields`	integer	Metadata lines recognised in the record
`has_body`	boolean	Record has body text

Column	Type	Content
`body_words`, `body_chars`	integer	Word and character counts of the body
`body_sha256`	string	SHA-256 hash of the body text
`text`	string	Article text as exported (`Texto completo`); empty if the record has none

3.2 Companion files

- `data/interim/exclusions.parquet`: `source_file`, `record_index` and one exclusion `reason` per row of the parsed table, footers and repeats included (section 6).
- `reports/tables/raw_by_year.csv` and `reports/tables/raw_by_download_date.csv`: records and distinct ids per year, and counts and publication-year range per download date, all as parsed, before repeated ids were removed.
- `reports/figures/records_by_year.svg` and `reports/figures/repeats_by_year.svg`: Figures 1 and 2.

4. Contents

4.1 Newspapers

With author: share of records with a named author.

Newspaper	Records	Print	Online	Pre-1997	Blog/podcast/website	With author
New York Times	387,610	148,077	239,533	0	239,533	33.1%
Washington Post	275,422	129,755	94,982	50,685	94,982	49.7%
Wall Street Journal	121,396	45,199	76,197	0	0	66.7%
Los Angeles Times	144,566	82,026	15,182	47,358	15,182	55.5%
Boston Globe	103,429	57,427	8,685	37,317	0	57.3%
Chicago Tribune	72,908	26,214	23,006	23,688	23,006	50.3%
USA Today	36,918	13,494	19,248	4,176	0	73.6%
Newsday	51,387	51,387	0	0	0	63.4%
Tampa Bay Times	12,921	9,985	2,936	0	2,936	22.5%
Star Tribune	7,001	7,001	0	0	0	77.3%
Other	12,413	11,870	530	13	797	54.8%
No metadata	885	0	0	0	0	0.0%
Total	1,226,856	582,435	480,299	163,237	376,436	48.7%

4.2 Online editions and the blog label

Newspaper	Online records	First year	Labelled blog/podcast/website	Years with both labels
New York Times	239,533	1996	100.0%	0
Washington Post	94,982	2011	100.0%	0
Wall Street Journal	76,197	2010	0.0%	0
Los Angeles Times	15,182	2017	100.0%	0
Boston Globe	8,685	2006	0.0%	0

Newspaper	Online records	First year	Labelled blog/podcast/website	Years with both labels
Chicago Tribune	23,006	2017	100.0%	0
USA Today	19,248	2014	0.0%	0
Tampa Bay Times	2,936	2017	100.0%	0

4.3 Blog-labelled records by document type

Document type	Blog-labelled records	Share of blog-labelled
News	187,420	49.8%
Opinions	142,548	37.9%
Commentary	52,706	14.0%
Editorial	49,769	13.2%
Letter To The Editor	242	0.1%
New s	116	0.0%

4.4 Years

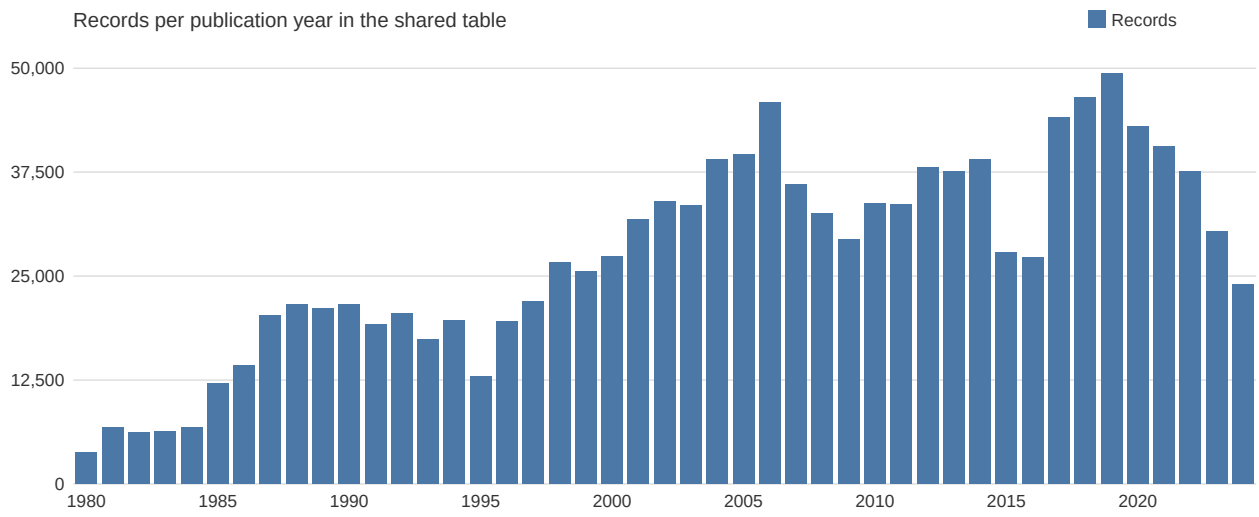


Figure 1. Records per publication year in the shared table.

Decade	Records	With author
1980s	119,028	47.0%
1990s	204,934	40.9%
2000s	349,144	36.4%
2010s	377,225	54.4%
2020s	175,619	71.2%
Unknown year	906	1.9%

4.5 Document types and authorship

A record with two labels counts once under each. Editorials are mostly unsigned; commentary and op-eds mostly signed.

Document type	Records	With author
News	334,226	43.6%
Editorial	275,482	24.0%
NEWSPAPER	266,015	44.6%
Opinions	236,229	58.2%
Commentary	212,130	88.5%
Op-Ed	8,835	94.6%
OP-ED	6,044	94.6%
EDITORIAL	5,603	1.4%
Opinion	5,119	87.3%
Letter To The Editor	2,677	0.1%
editorial	2,484	1.8%
Opinion Piece	2,287	98.3%

4.6 Sections

Top labels, grouped ignoring case. Each newspaper names its sections differently.

Section	Records	Share
opinion	383,596	31.3%
editorial	112,459	9.2%
a	108,337	8.8%
editorial-opinion	76,542	6.2%
opinions	52,325	4.3%
op/ed	36,991	3.0%
news	32,570	2.7%
editorial page	22,914	1.9%
editorial - opinion	21,099	1.7%
op-ed	17,909	1.5%
california metro; part b; editorial pages desk	16,498	1.3%
main news; part a; editorial desk	12,490	1.0%

4.7 Language

Language	Records	Share
English only	1,224,710	99.8%
Includes Spanish	1,261	0.1%
Missing	885	0.1%

4.8 Missing values

Field	Missing	Share
Title	888	0.1%

Field	Missing	Share
Author	629,809	51.3%
Section	48,953	4.0%
Full publication date	20,516	1.7%
Document id	882	0.1%
Body text	1,622	0.1%

4.9 Publications outside the query

Most common first. These are removed in the cleaned dataset (reason `not_queried_newspaper`).

Publication	Records
La Opinión; Los Angeles, Calif.	1,247
Washington Examiner; Washington, D.C.	1,204
The Press Democrat; Santa Rosa, Calif.	396
Philadelphia Daily News; Philadelphia, Pa.	377
South Florida Sun - Sentinel; Fort Lauderdale, Fla.	299
Pittsburgh Post - Gazette; Pittsburgh, Pa.	237
Politico, U.S. edition; Arlington	192
Orlando Sentinel (Online); Orlando	187
Public Opinion; Chambersburg, Pa.	185
Daily News; Los Angeles, Calif.	177
Reno Gazette - Journal; Reno, Nev.	161
York Daily Record; York, Pa.	158

5. Technical validation

5.1 Do the filters hold?

Filter	Result	Section
Newspapers	296 publications outside the list also appear. Queried newspapers with no records: none	4.1, 4.9
Source type	Not newspapers only, by design: 376,436 records (30.7%) are online editions that ProQuest labels `Blog, podcast o sitio web`	4.1–4.3
Document type	117 distinct labels; 402,716 records (32.8%) carry none of Commentary, Editorial or News	4.5
English	99.8% English only; the rest bilingual or unlabelled	4.7
Dates	1980–2024	4.4

The blog label:

- 99.9% of blog-labelled records are `(Online)` editions, mostly New York Times, Washington Post, Los Angeles Times, Chicago Tribune, Tampa Bay Times.
- It belongs to the publication edition, not the article or date: per newspaper, all online records carry it or none do, in every year (section 4.2, last column).
- 95.1% sit in an opinion or editorial section (section 4.3). The content is news opinion, not casual blog posts.

5.2 Repeats

Of 1,374,965 parsed records, 145,118 (10.6%) repeat a ProQuest id; 1,228,965 ids are distinct. 100,661 ids appear in more than one file and 0 repeat inside a file, which points to overlapping downloads. 2014 stands out in Figure 2: 112,098 records but only 39,087 distinct ids, because its articles were downloaded repeatedly (69.1% of its ids appear in 3 or more files). The peak is duplication, not more articles.

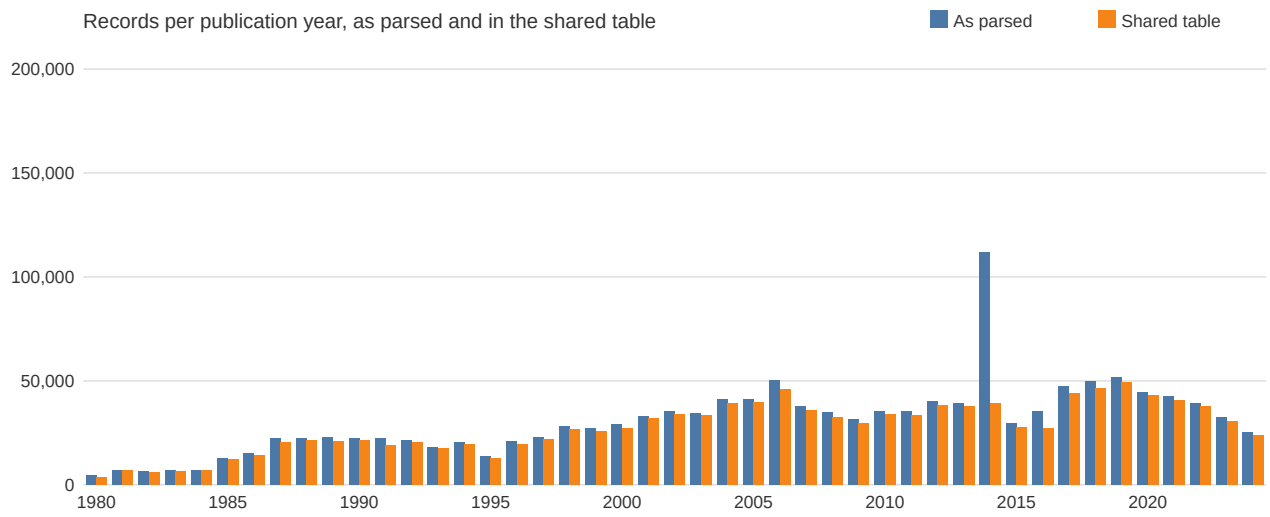


Figure 2. Records per publication year as parsed and in the shared table. The gap is the repeats removed.

Apart from stray spaces, the copies of an id agree on date, newspaper, author, section and document type. In 878 ids one copy has no body next to a copy with text, so the copy kept is the one with the longest body, and the earliest download if they tie. The others are removed from the shared table.

A second kind of repeat has another id but the same body, title, date and queried newspaper (editions of one newspaper count as one): 2,991 records. The copy kept is the earliest one that qualifies as a signed opinion piece, or the earliest of all if none does; the others are removed. Publications outside the query are judged only as not queried, so repeats among them stay.

5.3 News and opinion

Document type is not a reliable marker of opinion. 36.3% of records have a non-opinion label but sit in an opinion section, and 150,950 (12.3%) have neither an opinion label nor an opinion section. All are kept in the raw table.

6. Usage notes and limitations

- **Exclusion reasons:** the companion table gives every parsed row one reason: file footer, repeated id, publication not queried, repeated text (same body, title, date and newspaper), missing metadata, not opinion, unclear opinion, no text, no author, no person among the authors, or kept. Only file footer, repeated id and repeated text leave the shared table; the others say which records the cleaned dataset leaves out and do not change this table. Counts: `reports/tables/exclusion_counts.csv`.
- **Subtler duplicates** (reprints, edited versions under another title) are not detected.
- **Authors:** 597,047 records (48.7%) name an author. Unsigned pieces, mostly editorials, cannot be linked to a person.
- **No body text:** 1,622 records, all in 5 files.

- **Metadata as exported:** values keep ProQuest's stray spaces; use the `*_key` columns to compare categories. Fields that are sparse (for example `url`, `company`, `publication_history`) are kept in case a later stage needs them.
- **License:** the table contains ProQuest text and is for internal use among the project's researchers only.
- **Query matches text, not genre:** the keyword search returns any article mentioning "opinion" or "editorial", and News is a queried document type.

7. Code availability

The pipeline is in the private repository `econ_corpus` [2]: `parse` (exports to table), `export` (table to CSV), `reasons` and `exclusions`, `tables`, `figures` and `datasheet` (this document). Every rule and its evidence are in `docs/decisions.md`. `README.md` lists the commands, from `make parse` on, that regenerate the table and this document. The repository holds code and documentation only; the data are not in it.

8. References

[1] Statista, "Largest daily newspapers in the United States in the six months to September 2023, by average print circulation (in thousands)," Statista, Feb. 19, 2024. Accessed: Nov. 3, 2026. [Online]. Available: https://www.statista.com/statistics/272790/circulation-of-the-biggest-daily-newspapers-in-the-us/

[2] J. C. Reyes, "econ_corpus: build and audit a corpus of U.S. newspaper opinion articles," GitHub repository, 2026. Accessed: Nov. 3, 2026. [Online]. Available: https://github.com/JuanReyes01/econ_corpus